

Relational Algebra

Operations in the Relational Model

These operation can be expressed in an algebra, called **“relational algebra.”**

In this algebra, **relations** are the **operands** and we apply **operators** on them.

Operations

Four broad classes:

1. Usual set operations

union

intersection

difference

2. Operations that remove parts of a relation:

selection eliminates some rows(tuples)

projection eliminates some columns

3. Operations that combine the tuples of two relations:

Cartesian product pairs tuples of two relations in all possible ways

join selectively pairs tuples from two relations.

4. An operation called “**renaming**”.

Conditions for Set Operations on Relations

We can apply **union**, **intersection**, **difference**- on relations **R** and **S** provided that:

1. **R** and **S** must have schemas with identical sets of attributes.
2. Before applying the operations, the columns of **R** and **S** must be ordered so that the order of attributes is the same for both relations.

Set Operations on Relations

$R \cup S$, the **union** of R and S , is the set of tuples that are in R or S or both.

$R \cap S$, the **intersection** of R and S , is the set of tuples that are in both R and S .

$R - S$, the **difference** of R and S , is the set of tuples that are in R but not in S .

Note that $R - S$ is different from $S - R$.

Projection

$\pi_{A_1, \dots, A_n}(\mathbf{R})$

Produces from relation \mathbf{R} a new relation that has only the A_1, \dots, A_n columns of \mathbf{R} .

Example:

For $\pi_{\text{title, year, length}}(\mathbf{Movies})$ on:

| <i>title</i> | <i>year</i> | <i>length</i> | <i>filmType</i> | <i>studioName</i> | <i>producerC#</i> |
|---------------|-------------|---------------|-----------------|-------------------|-------------------|
| Star wars | 1977 | 124 | color | Fox | 12345 |
| Mighty Ducks | 1991 | 104 | color | Disney | 67890 |
| Wayne's World | 1992 | 95 | color | Paramount | 99999 |

Example (Continued)

We get:

| <i>title</i> | <i>year</i> | <i>length</i> |
|---------------|-------------|---------------|
| Star wars | 1977 | 124 |
| Mighty Ducks | 1991 | 104 |
| Wayne's World | 1992 | 95 |

What about π_{filmtype} (**Movies**) ?

Selection

$\sigma_c(\mathbf{R})$

Produces a new relation with those tuples of \mathbf{R} which satisfy condition \mathbf{C} .

Example:

For $\sigma_{length \geq 100}(\mathbf{Movies})$ we have as result:

| <i>title</i> | <i>year</i> | <i>length</i> | <i>filmType</i> | <i>studioName</i> | <i>producerC#</i> |
|--------------|-------------|---------------|-----------------|-------------------|-------------------|
| Star wars | 1977 | 124 | color | Fox | 12345 |
| Mighty Ducks | 1991 | 104 | color | Disney | 67890 |

Another Example

Suppose we want the movies by Fox which are at least 100 minutes long.

$\sigma_{length \geq 100 \text{ And } studioName = 'Fox'}$ **(Movies)**

Result is:

| <i>title</i> | <i>year</i> | <i>length</i> | <i>filmType</i> | <i>studioName</i> | <i>producerC#</i> |
|--------------|-------------|---------------|-----------------|-------------------|-------------------|
| Star wars | 1977 | 124 | color | Fox | 12345 |

Cartesian Product

R×S

1. Set of tuples ***rs*** that are formed by choosing the first part (***r***) to be any tuple of **R** and the second part (***s***) to be any tuple of **S**.
2. Schema for the resulting relation is the union of schemas for **R** and **S**.
3. If **R** and **S** happen to have some attributes in common, then prefix those attributes by the relation name.

Example:

| | | | | | | |
|---|---|---|---|---|----|----|
| R | A | B | S | B | C | D |
| | 1 | 2 | | 2 | 5 | 6 |
| | 3 | 4 | | 4 | 7 | 8 |
| | | | | 9 | 10 | 11 |

Example (Continued)

Resulting relation will be:

| R×S | A | R.B | S.B | C | D |
|------------|----------|------------|------------|----------|----------|
| | 1 | 2 | 2 | 5 | 6 |
| | 1 | 2 | 4 | 7 | 8 |
| | 1 | 2 | 9 | 10 | 11 |
| | 3 | 4 | 2 | 5 | 6 |
| | 3 | 4 | 4 | 7 | 8 |
| | 3 | 4 | 9 | 10 | 11 |

Theta-Join

$\mathbf{R} \bowtie_{\mathbf{C}} \mathbf{S}$.

1. The result of this operation is constructed as follows:
 - a) Take the Cartesian product of **R** and **S**.
 - b) Select from the product only those tuples that satisfy the condition **C**.
2. Schema for the result is the union of the schema of **R** and **S** with, “**R**” or “**S**” prefix as necessary.

When the condition is equality, we call it “equijoin”.

Natural Join

Special case of equijoin when attributes we want to use in join have the same name in both tables

$\mathbf{R} \bowtie \mathbf{S}$

Let $\mathbf{A}_1, \mathbf{A}_2, \dots, \mathbf{A}_n$ be the attributes in both the schema of \mathbf{R} and the schema of \mathbf{S} .

Then a tuple \mathbf{r} from \mathbf{R} and a tuple \mathbf{s} from \mathbf{S} are successfully paired if and only if \mathbf{r} and \mathbf{s} agree on each of the attributes $\mathbf{A}_1, \mathbf{A}_2, \dots, \mathbf{A}_n$.

Example: The natural join of the relation \mathbf{R} and \mathbf{S} from previous example is:

| A | B | C | D |
|---|---|---|---|
| 1 | 2 | 5 | 6 |
| 3 | 4 | 7 | 8 |

Attributes with the same name have only one representative.

Why?

Example

Compute the natural and theta join for relations U and V:

| A | B | C |
|---|---|---|
| 1 | 2 | 3 |
| 6 | 7 | 8 |
| 9 | 7 | 8 |

Relation U

| B | C | D |
|---|---|----|
| 2 | 3 | 4 |
| 2 | 3 | 5 |
| 7 | 8 | 10 |

Relation V

$U \bowtie V$ and $U \bowtie_{A < D} V$

Example

| A | B | C | D |
|---|---|---|----|
| 1 | 2 | 3 | 4 |
| 1 | 2 | 3 | 5 |
| 6 | 7 | 8 | 10 |
| 9 | 7 | 8 | 10 |

Result $U \bowtie V$

| A | U.B | U.C | V.B | V.C | D |
|---|-----|-----|-----|-----|----|
| 1 | 2 | 3 | 2 | 3 | 4 |
| 1 | 2 | 3 | 2 | 3 | 5 |
| 1 | 2 | 3 | 7 | 8 | 10 |
| 6 | 7 | 8 | 7 | 8 | 10 |
| 9 | 7 | 8 | 7 | 8 | 10 |

Result of $U \bowtie_{A < D} V$

Combing Operations to Form Queries

“What are the **title** and **years** of movies made by **Fox** that are at least 100 minutes long?”

$\pi_{title, year}(\sigma_{length \geq 100 \text{ AND } studioName = 'Fox'}(Movies))$

```
SELECT title, year  
FROM Movies  
WHERE length >= 100 AND studioName = 'Fox';
```


Another Example

Consider two relations **Movies** and **StarsIn**,

With schemas:

Movies(title, year, length, filmType, studioName)

StarsIn(title, year, starName)

Suppose we want to know:

“Find the stars of the movies that are at least 100 minutes long.”

First we join the two relations: **Movies**, **StarsIn**

Second we select movies with length at least 100 min.

Then we project onto **starName**.

Another Example – Solution 1

$\pi_{\text{starName}}(\sigma_{\text{Movies.title=StarsIn.title AND Movies.year=StarsIn.year AND length}\geq 100}(\text{Movies} \times \text{StarsIn}))$

SELECT starName

FROM Movies, StarsIn

WHERE Movies.title=StarsIn.title AND Movies.year=StarsIn.year

AND length>=100;

Another Example – Solution 2

$\pi_{\text{starName}}(\sigma_{\text{length} \geq 100}(\text{Movies} \bowtie_{\text{Movies.title=StarsIn.title AND Movies.year=StarsIn.year}} \text{StarsIn}))$

SELECT starName

FROM Movies JOIN StarsIn ON Movies.title=StarsIn.title AND
Movies.year=StarsIn.year

WHERE length >= 100;

Another Example – Solution 3

$\pi_{\text{starName}}(\sigma_{\text{length} \geq 100}(\text{Movies} \bowtie \text{StarsIn}))$

```
SELECT starName
FROM Movies JOIN StarsIn USING (title, year)
WHERE length >= 100;
```

Or (less safe):

```
SELECT starName
FROM Movies NATURAL JOIN StarsIn
WHERE length >= 100;
```

Renaming Operator

$\rho_{S(A_1, A_2, \dots, A_n)}(\mathbf{R})$

1. Resulting relation has exactly the same tuples as \mathbf{R} , but the name of the relation is \mathbf{S} .
2. Moreover, the attributes of the result relation \mathbf{S} are named A_1, A_2, \dots, A_n , in order from the left.

Problem

Product(maker, model, type)

PC(model, speed, ram, hd, rd, price)

Laptop(model, speed, ram, hd, screen, price)

Printer(model, color, type, price)

(Exercise 5.2.1)

| <i>maker</i> | <i>model</i> | <i>type</i> |
|--------------|--------------|-------------|
| A | 1001 | pc |
| A | 1002 | pc |
| A | 1003 | pc |
| A | 2004 | laptop |
| A | 2005 | laptop |
| A | 2006 | laptop |
| B | 1004 | pc |
| B | 1005 | pc |
| B | 1006 | pc |
| B | 2001 | laptop |
| B | 2002 | laptop |
| B | 2003 | laptop |
| C | 1007 | pc |
| C | 1008 | pc |
| C | 2008 | laptop |
| C | 2009 | laptop |
| C | 3002 | printer |
| C | 3003 | printer |
| C | 3006 | printer |
| D | 1009 | pc |
| D | 1010 | pc |
| D | 1011 | pc |
| D | 2007 | laptop |
| E | 1012 | pc |
| E | 1013 | pc |
| E | 2010 | laptop |
| F | 3001 | printer |
| F | 3004 | printer |
| G | 3005 | printer |
| H | 3007 | printer |

| <i>model</i> | <i>speed</i> | <i>ram</i> | <i>hd</i> | <i>rd</i> | <i>price</i> |
|--------------|--------------|------------|-----------|-----------|--------------|
| 1001 | 700 | 64 | 10 | 48xCD | 799 |
| 1002 | 1500 | 128 | 60 | 12xDVD | 2499 |
| 1003 | 866 | 128 | 20 | 8xDVD | 1999 |
| 1004 | 866 | 64 | 10 | 12xDVD | 999 |
| 1005 | 1000 | 128 | 20 | 12xDVD | 1499 |
| 1006 | 1300 | 256 | 40 | 16xDVD | 2119 |
| 1007 | 1400 | 128 | 80 | 12xDVD | 2299 |
| 1008 | 700 | 64 | 30 | 24xCD | 999 |
| 1009 | 1200 | 128 | 80 | 16xDVD | 1699 |
| 1010 | 750 | 64 | 30 | 40xCD | 699 |
| 1011 | 1100 | 128 | 60 | 16xDVD | 1299 |
| 1012 | 350 | 64 | 7 | 48xCD | 799 |
| 1013 | 733 | 256 | 60 | 12xDVD | 2499 |

(a) Sample data for relation PC

| <i>model</i> | <i>speed</i> | <i>ram</i> | <i>hd</i> | <i>screen</i> | <i>price</i> |
|--------------|--------------|------------|-----------|---------------|--------------|
| 2001 | 700 | 64 | 5 | 12.1 | 1448 |
| 2002 | 800 | 96 | 10 | 15.1 | 2584 |
| 2003 | 850 | 64 | 10 | 15.1 | 2738 |
| 2004 | 550 | 32 | 5 | 12.1 | 999 |
| 2005 | 600 | 64 | 6 | 12.1 | 2399 |
| 2006 | 800 | 96 | 20 | 15.7 | 2999 |
| 2007 | 850 | 128 | 20 | 15.0 | 3099 |
| 2008 | 650 | 64 | 10 | 12.1 | 1249 |
| 2009 | 750 | 256 | 20 | 15.1 | 2599 |
| 2010 | 366 | 64 | 10 | 12.1 | 1499 |

(b) Sample data for relation Laptop

| <i>model</i> | <i>color</i> | <i>type</i> | <i>price</i> |
|--------------|--------------|-------------|--------------|
| 3001 | true | ink-jet | 231 |
| 3002 | true | ink-jet | 267 |
| 3003 | false | laser | 390 |
| 3004 | true | ink-jet | 439 |
| 3005 | true | bubble | 200 |
| 3006 | true | laser | 1999 |
| 3007 | false | laser | 350 |

(c) Sample data for relation Printer

Figure 5.10: Sample data for Product

Problem

- a) Which PC models have a speed of at least 1000?
- b) Which manufacturers make laptops with a hard disk of at least 30?
- c) Find the model number and price of all products (of any type) made by manufacturer *B*.
- d) Find the model numbers of all color laser printers.
- e) Find those manufacturers that sell Laptops, but not PC's.
- !f) Find those hard-disk sizes that occur in two or more PC's.
- !g) Find those pairs of PC models that have both the same speed and RAM. A pair should be listed once.
- !!h) Find those manufacturers of at least two different computers (PC or Laptops) with speed of at least 700.