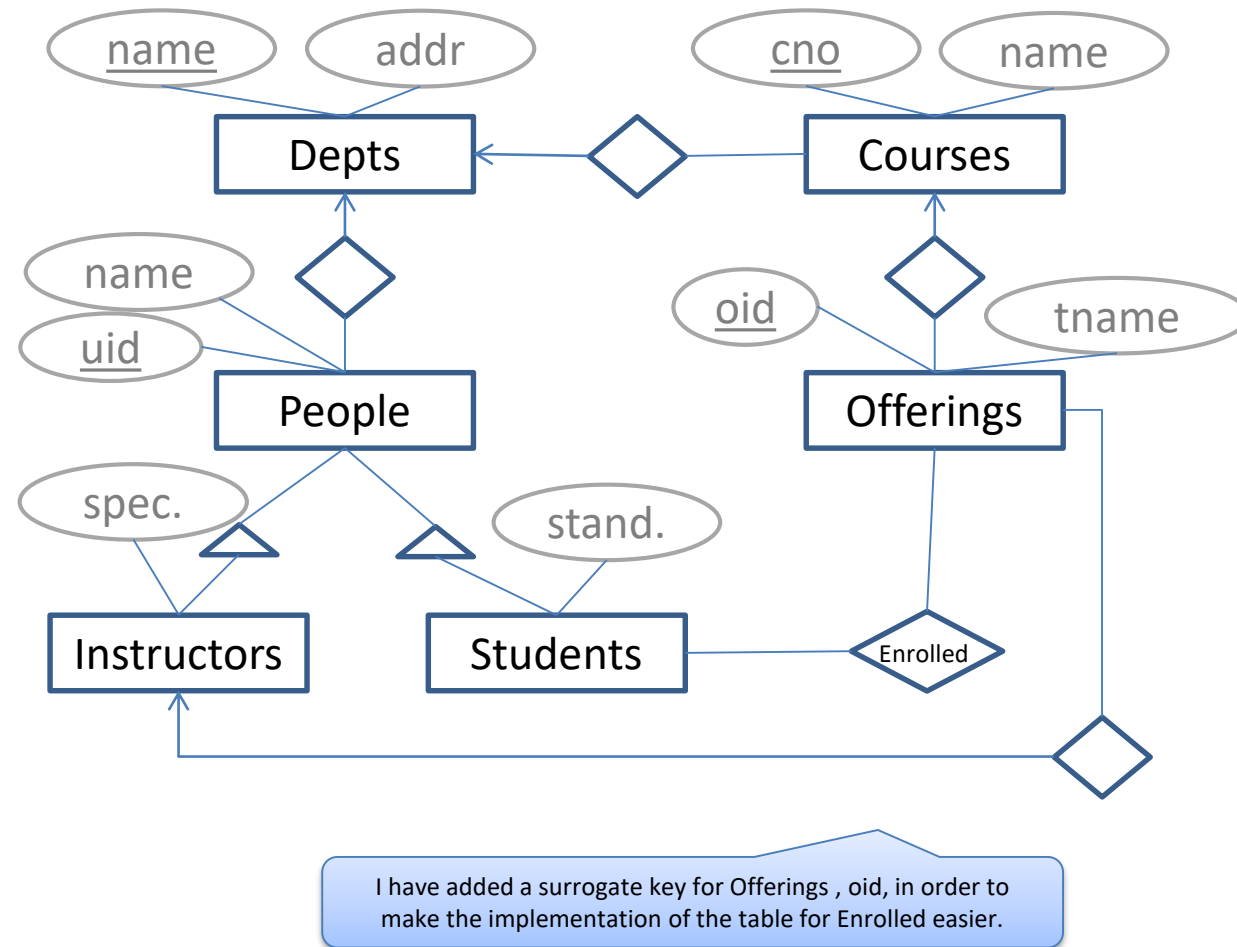


# Review

# ER, Tables, Queries, Constraints



```
CREATE TABLE Depts (
  name VARCHAR(20) PRIMARY KEY,
  addr VARCHAR(20)
);
```

```
CREATE TABLE People (
  uid CHAR(10) PRIMARY KEY,
  name VARCHAR(20),
  dname VARCHAR(20) REFERENCES Depts(name)
);
```

```
CREATE TABLE Instructors (
  uid CHAR(10) PRIMARY KEY REFERENCES People(uid),
  specialty VARCHAR(40)
);
```

```
CREATE TABLE Students (
  uid CHAR(10) PRIMARY KEY REFERENCES People(uid),
  standing VARCHAR(30)
);
```

```
CREATE TABLE Courses (
  cno CHAR(8) PRIMARY KEY,
  name VARCHAR(40),
  dname VARCHAR(20) REFERENCES Depts(name)
);
```

```
CREATE TABLE Offerings (
  oid INT PRIMARY KEY,
  cno CHAR(8) REFERENCES Courses(cno),
  tname CHAR(6),
  uid CHAR(10) REFERENCES Instructors(uid)
);
```

```
CREATE TABLE Enrolled (
  oid INT REFERENCES Offerings(oid),
  uid CHAR(10) REFERENCES Students(uid),
  PRIMARY KEY(oid,uid)
);
```

```
INSERT INTO Depts VALUES('Comp.Sci.', 'ECS Bldg');
INSERT INTO People VALUES('V11111111', 'Jon', 'Comp.Sci.');
INSERT INTO People VALUES('V22222222', 'Ben', 'Comp.Sci.');
INSERT INTO Instructors VALUES('V11111111', 'Hardware');
INSERT INTO Students VALUES('V22222222', '3rd year');
INSERT INTO Courses VALUES('CSC390', 'Hardware Systems', 'Comp.Sci.');
INSERT INTO Offerings VALUES(1, 'CSC390', '201409', 'V11111111');
INSERT INTO Enrolled VALUES(1, 'V22222222');
```

# Mutually exclusive subclasses

```
CREATE TABLE Vehicles (  
    vin CHAR(17) PRIMARY KEY,  
    vehicle_type CHAR(3) CHECK(vehicle_type IN ('SUV', 'ATV')),  
    fuel_type CHAR(4),  
    door_count INT CHECK(door_count >= 0),  
    UNIQUE(vin, vehicle_type)  
);  
  
CREATE TABLE SUVs (  
    vin CHAR(17) PRIMARY KEY,  
    vehicle_type CHAR(3) CHECK(vehicle_type = 'SUV'),  
    FOREIGN KEY (vin, vehicle_type) REFERENCES Vehicles (vin, vehicle_type)  
        ON DELETE CASCADE  
);  
  
CREATE TABLE ATVs (  
    vin CHAR(17) PRIMARY KEY,  
    vehicle_type CHAR(3) CHECK(vehicle_type = 'ATV'),  
    FOREIGN KEY (vin, vehicle_type) REFERENCES Vehicles (vin, vehicle_type)  
        ON DELETE CASCADE  
);
```

# Views with check option: Example

```
CREATE TABLE Hotel (  
    room_nbr INT NOT NULL,  
    arrival_date DATE NOT NULL,  
    departure_date DATE NOT NULL,  
    guest_name CHAR(15) NOT NULL,  
    PRIMARY KEY (room_nbr, arrival_date),  
    CHECK (departure_date > arrival_date)  
);
```

We want to add the constraint that reservations do not overlap.

```
CREATE VIEW HotelStays AS  
SELECT room_nbr, arrival_date, departure_date, guest_name  
FROM Hotel H1  
WHERE NOT EXISTS (  
    SELECT *  
    FROM Hotel H2  
    WHERE H1.room_nbr = H2.room_nbr AND  
        (H2.arrival_date < H1.arrival_date AND H1.arrival_date < H2.departure_date)  
)  
WITH CHECK OPTION;
```

# SQL

SELECT ... FROM ... WHERE ... GROUP BY ... HAVING ... ORDER BY

JOIN ON

JOIN USING

LEFT OUTER JOIN

RIGHT OUTER JOIN

LIKE

IS NULL

IS NOT NULL

Careful what can go to SELECT when having GROUP BY

Careful regarding the difference in conditions that go to HAVING vs those that go to WHERE.

# Example

- Using **Movies**, **StarsIn**, and **Stars**,  
*find the star's total length of film played.*  
We are interested only in Canadian stars and  
who first appeared in a movie before 2000.

```
SELECT starName, SUM(length)
FROM Movies, StarsIn, Stars
WHERE Movies.title=StarsIn.title AND Movies.year=StarsIn.year
      AND Stars.name=StarsIn.starName
      AND Stars.birthplace LIKE '%Canada%'
GROUP BY starName
HAVING MIN(StarsIn.year) < 2000;
```

# Correlated Subqueries

- Suppose StarsIn table has an additional attribute “salary”

**StarsIn**(movie, movie, starName, salary)

Now, find the stars who were paid for some movie more than the average salary for that movie.

```
SELECT starName, title, year
FROM StarsIn X
WHERE salary >
      (SELECT AVG(salary)
       FROM StarsIn
       WHERE title = X.title AND year=X.year);
```

## Remark

Semantically, the value of the X tuple changes in the outer query, so the database must rerun the subquery for each X tuple.



# Another Solution (Nesting in FROM)

```
SELECT X.starName, X.title, X.year
FROM StarsIn X, (SELECT title, year, AVG(salary) AS avgSalary
                  FROM StarsIn
                  GROUP BY title, year) Y
WHERE X.salary>Y.avgSalary AND
      X.title=Y.title AND X.year=Y.year;
```

# Intersection, Union, Difference


```
SELECT title, year
FROM StarsIn
WHERE starName='Richard Gere'
    INTERSECT
SELECT title, year
FROM StarsIn
WHERE starName='Julia Roberts' ;
```

```
SELECT title, year
FROM StarsIn
WHERE starName='Richard Gere'
    UNION
SELECT title, year
FROM StarsIn
WHERE starName='Julia Roberts' ;
```

```
SELECT title, year
FROM StarsIn
WHERE starName='Richard Gere'
    EXCEPT
SELECT title, year
FROM StarsIn
WHERE starName='Julia Roberts' ;
```

# Storage

# Disks



- **Platters** with top and bottom surfaces rotate around a spindle.
- **2--30 surfaces.**
- **Rotation speed:** 3600--7200 rpm.
- **One head per surface.**
- **All heads move in and out in unison.**

# **AVG** time to read a **16,384**-byte block (Megatron 747 – a fictitious, but realistic disk...)


- Transfer time is **0.13** milliseconds and
- Average rotational latency is the time to rotate the disk half way around, or **4.17** milliseconds.
- Average seek time is:  $1 + (65536/3) * (1/4000) = \mathbf{6.46}$  ms
- Total:  $6.46 + 4.17 + 0.13 = \mathbf{10.76}$  ms (or about **11 ms**)

# 2PMMS

## Phase 1

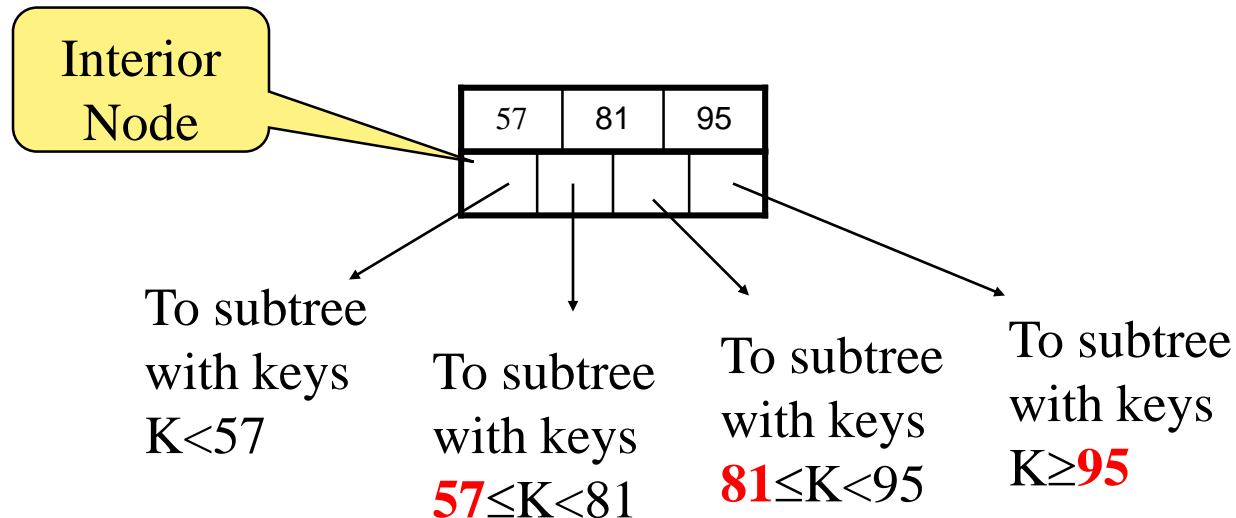

1. Fill main memory with records.
2. Sort using favorite main memory sort.
3. Write sorted sublist to disk.
4. Repeat until all records have been put into one of the sorted lists.

## Phase 2



# Indexes

# BTrees: A typical leaf and interior node (unclustered index)





**57, 81, and 95** are the least keys we can reach by via the corresponding pointers.



# A typical leaf and interior node (clustered index)

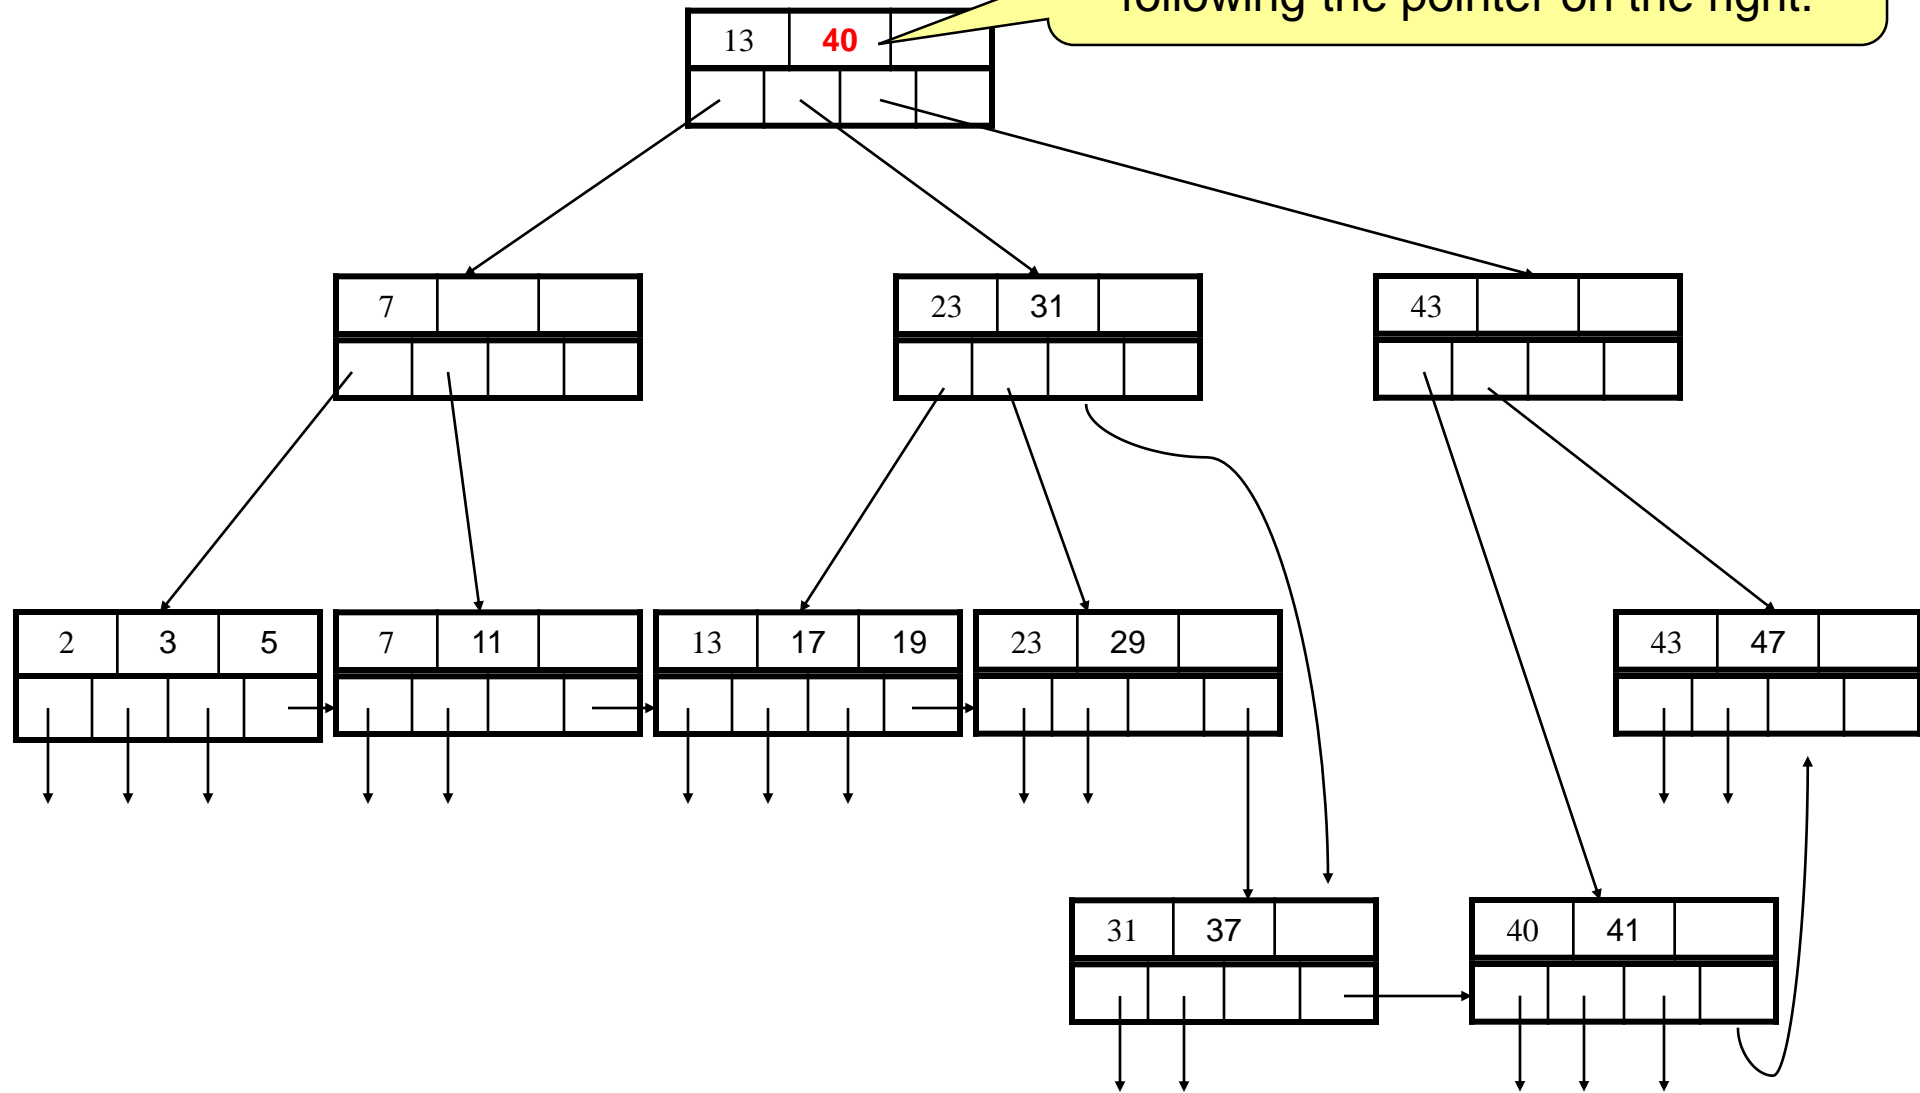
A clustered index is by necessity primary.



**57, 81, and 95** are the least keys we can reach by via the corresponding pointers.

# Insertion of key 40

40 is the least key reachable by following the pointer on the right.



# Structure of B-trees with real blocks

- Degree  $n$  means that all nodes have space for  $n$  search keys and  $n+1$  pointers
- **Node = block**
- Let
  - block size be 16,384 Bytes,
  - key 20 Bytes,
  - pointer 20 Bytes.
- Let's solve for  $n$ :

$$20n + 20(n+1) \leq 16384$$

$$\Rightarrow n \leq 409$$

# Query Evaluation

# An SQL query and its RA equiv.

**Employees** (sin INT, ename VARCHAR(20), rating INT, address VARCHAR(90))

**Maintenances** (sin INT, planeId INT, day DATE, desc CHAR(120))


SELECT ename

FROM Employees NATURAL JOIN Maintenances

WHERE planeId = 100 AND rating > 5;

$\pi_{ename}(\sigma_{planeId=100 \text{ AND } rating>5}(\text{Employees} \bowtie \text{Maintenances}))$

RA expressions can be represented by an expression tree.



An algorithm is chosen for each node in the expression tree.

**Initial trees** can be transformed to "better" trees.

Process of finding **good evaluation plans** is called **query optimization**.

# Running Example – Airline

**Employees** (sin INT, ename VARCHAR(20), rating INT, address VARCHAR(90))

**Maintenances** (sin INT, planeId INT, day DATE, desc CHAR(120))

- Assume for **Maintenances**:
  - a tuple is **160 bytes**
  - a block can hold **100 tuples** (16K block)
  - we have **1000 blocks** of such tuples.
- Assume for **Employees**:
  - a tuple is **130 bytes**
  - a block can hold **120 tuples**
  - we have **50 blocks** of such tuples.

# Index nested loops join


- Scan **Maintenances** and for each tuple **probe** **Employees** for matching tuples (using the index on **Employees.SIN**).
- **Analysis:**
  - For each of the 100,000 maintenance tuples, we try to access the corresponding employee with 3 I/Os (via the index).
  - So,  $100,000 * 3 = 300,000$  I/Os !!

# Sort-merge join

- **Sort** both tables **on the join column**, then scan them to find matches
- **Analysis:**
  - Sort **Maintenances** with 2PMMS, and **Employees** with 2PMMS
  - Cost for sort is
    - $4 * 1000 = 4000$  I/Os for **Maintenances** and
    - $4 * 50 = 200$  I/Os. for **Employees**
  - Then we merge-join. This requires an additional scan of both tables.
  - Thus the total cost is  $4000 + 200 + 1000 + 50 = \mathbf{5250}$  I/Os. (**Much better!!**)
- However, “**index nested loops**” method has the nice property that it is **incremental**.



# Cost of a Plan (Pushing selections)



$$(1000+10) + (50+25) + (4*10 + 4*25) + (10+25) = \mathbf{1260} \text{ I/Os}$$

1000 I/Os to scan  
Maintenances,  
10 I/Os to save  
temporary table T1  
(assuming there  
are 100 planes and  
for each we have  
an equal number of  
maintenance  
records)

50 I/Os to scan  
Employees,  
25 I/Os to save  
temporary table T2  
(assuming there  
are 10 ratings  
uniformly  
distributed)

Cost for sorting T1  
and T2 on SIN  
(the join attribute)  
using 2PMMS.

Cost for reading  
the sorted T1 and  
T2 and doing the  
join.

See the full slides  
more examples

# Transactions

Controlling Concurrent Behavior

# Summarizing the Terminology

- **Transaction** (model) is a *sequence* of  $r$  and  $w$  actions on database elements.
- **Schedule** is a *sequence* of read/write actions performed by a collection of transactions.
- **Serial Schedule** = All actions for each transaction are consecutive.  
 $r_1(A); w_1(A); r_1(B); w_1(B); r_2(A); w_2(A); r_2(B); w_2(B); \dots$
- **Serializable Schedule**: A schedule whose “**effect**” is equivalent to that of some serial schedule.
- **Conflict-serializable Schedule** if it can be converted into a serializable schedule with the **same effect** by a series of non-conflicting swaps of adjacent elements

We talk about a **sufficient condition** for serializability.

# Example of a conflict-serializable schedule

$r_1(A); w_1(A); r_2(A); w_2(A); r_1(B); w_1(B); r_2(B); w_2(B)$

$r_1(A); w_1(A); r_2(A); \underline{w_2(A)}; \underline{r_1(B)}; w_1(B); r_2(B); w_2(B)$   
 $r_1(A); w_1(A); \underline{r_2(A)}; \underline{r_1(B)}; w_2(A); w_1(B); r_2(B); w_2(B)$   
 $r_1(A); w_1(A); r_1(B); r_2(A); \underline{w_2(A)}; \underline{w_1(B)}; r_2(B); w_2(B)$   
 $r_1(A); w_1(A); r_1(B); \underline{r_2(A)}; \underline{w_1(B)}; w_2(A); r_2(B); w_2(B)$   
 $r_1(A); w_1(A); r_1(B); w_1(B); r_2(A)w_2(A); r_2(B); w_2(B)$

The operations in bold can be safely swapped.

# Precedence graphs


$r_2(A); r_1(B); w_2(A); r_3(A); w_1(B); w_3(A); r_2(B); w_2(B)$

Note the following:

- $w_1(B) <_s r_2(B)$
- $r_2(A) <_s w_3(A)$

➤ These are conflicts since they contain a read/write on the same element

➤ They cannot be swapped. Therefore  $T_1 < T_2 < T_3$




Conflict serializable

$r_2(A); r_1(B); w_2(A); r_2(B); r_3(A); w_1(B); w_3(A); w_2(B)$

Note the following:

- $r_2(B) <_s w_1(B)$
- $w_2(A) <_s w_3(A)$
- $r_1(B) <_s w_2(B)$

➤ Here, we have  $T_1 < T_2 < T_3$ , but we also have  $T_2 < T_1$



Not conflict serializable

# Legal Schedule Doesn't Mean Serializable

A **scheduler** takes requests from transactions for reads and writes, and decides if it is "OK" to allow them to operate on DB or defer them until it is safe to do so.

One approach is to use **locks**.

T <sub>1</sub>	T <sub>2</sub>	A	B
		25	25
l <sub>1</sub> (A); r <sub>1</sub> (A)			
A = A + 100			
w <sub>1</sub> (A); u <sub>1</sub> (A)		125	
	l <sub>2</sub> (A); r <sub>2</sub> (A)		
	A = A * 2		
	w <sub>2</sub> (A); u <sub>2</sub> (A)	250	
	l <sub>2</sub> (B); r <sub>2</sub> (B)		
	B = B * 2		
	w <sub>2</sub> (B); u <sub>2</sub> (B)		50
l <sub>1</sub> (B); r <sub>1</sub> (B)			
B = B + 100			
w <sub>1</sub> (B); u <sub>1</sub> (B)			150

Consistency constraint assumed for this example:  
A=B

# Two Phase Locking

There is a simple condition, which guarantees conflict-serializability:

In every transaction, all lock requests (phase 1) precede all unlock requests (phase 2).

T <sub>1</sub>	T <sub>2</sub>	A	B
		25	25
l <sub>1</sub> (A); r <sub>1</sub> (A)			
A = A + 100			
w <sub>1</sub> (A); l <sub>1</sub> (B); u <sub>1</sub> (A)		125	
	l <sub>2</sub> (A); r <sub>2</sub> (A)		
	A = A * 2		
	w <sub>2</sub> (A)	250	
	l <sub>2</sub> (B) <b>Denied</b>		
r <sub>1</sub> (B)			
B = B + 100			125
w <sub>1</sub> (B); u <sub>1</sub> (B)			
	l <sub>2</sub> (B); u <sub>2</sub> (A); r <sub>2</sub> (B)		
	B = B * 2		
	w <sub>2</sub> (B); u <sub>2</sub> (B)		250

# Shared/Exclusive Locks

	S	X
S	yes	no
X	no	no



# Shared/Exclusive Locks Example

$r_1(A); r_2(B); r_3(C); r_1(B); r_2(C); r_3(D); w_1(A); w_2(B); w_3(C);$

$\underline{T_1}$  —————  $\underline{T_2}$  —————  $\underline{T_3}$   
 $xl_1(A); r_1(A)$

$xl_2(B); r_2(B)$

$xl_3(C); r_3(C)$

$sl_1(B)$  **denied**

$sl_2(C)$  **denied**

$sl_3(D); r_3(D);$   
 $w_3(C); u_3(C); u_3(D)$

$sl_2(C); r_2(C);$   
 $w_2(B);$   
 $u_2(B); u_2(C)$

$sl_1(B); r_1(B);$   
 $w_1(A);$   
 $u_1(A); u_1(B)$


If a transaction  $T_i$  only reads element  $X$ , it requests  $sl_i(X)$ .

However, if  $T_i$  reads and then writes  $X$ , it only requests  $xl_i(X)$ .

This example shows that when there are transactions that will eventually write the elements they read, having shared and exclusive locks isn't much better than having just simple locks.

# Upgrading Locks

$r_1(A); r_2(B); r_3(C); r_1(B); r_2(C); r_3(D); w_1(A); w_2(B); w_3(C);$



Now, with locks that can be upgraded, transactions ask first for shared locks to read elements, then ask for the locks to be upgraded to exclusive when they want to write those elements.

Everything goes smoothly without any request being denied.

# Possibility for Deadlocks

**Example:** T1 and T2 each reads X and later writes X.

$T_1$	$T_2$
$sl_1(X)$	
	$sl_2(X)$
$xl_1(X)$ Denied	
	$xl_2(X)$ Denied


**Problem:** when we allow upgrades, it is easy to get into a deadlock situation.

# Solution: Update Locks

- Update lock  $ul_i(X)$ .
  - Only an update lock (not shared lock) can be upgraded to exclusive lock (if there are no shared locks anymore).
  - A transaction that will read and later on write some element A, asks initially for an update lock on A, and then asks for an exclusive lock on A. Such transaction doesn't ask for a shared lock on A.

	S	X	U
S	yes	no	yes
X	no	no	no
U	no	no	no

# Benefits of Update Locks



T5 will wait for T1-T4 to release the lock on A, but will stop T6-T9 obtaining a lock on A as otherwise its updating of the lock to exclusive would be blocked for a long time (until all the other transactions release the lock on A).

# **Recovery from Crashes**

# Undo Logging

log all “important actions”

**<START T>** -- transaction **T** started.

**<T,X,Old<sub>x</sub>>** -- database element **X** was modified;  
it used to have the value **Old<sub>x</sub>**

**<COMMIT T>** -- transaction **T** has completed

**<ABORT T>** -- transaction **T** couldn't complete successfully.

**Two rules:**

**U1:** Log records for element **X** must be on disk **before** any change to **X** appears on disk.

**U2:** If a transaction **T** commits, then **<COMMIT T>** must be written to log in disk **after** all elements changed by **T** are written to disk.

Force the log to be saved to disk by executing **Flush Log**.

## Example:

Action	t	Buff A	Buff B	A in HD	B in HD	Log
Read(A,t)	8	8		8	8	<Start T>
t:=t*2	16	8		8	8	
Write(A,t)	16	16		8	8	<T,A,8>
Read(B,t)	8	16	8	8	8	
t:=t*2	16	16	8	8	8	
Write(B,t)	16	16	16	8	8	<T,B,8>
<b>Flush Log</b>						
Output(A)	16	16	16	16	8	
Output(B)	16	16	16	16	16	<Commit T>
<b>Flush Log</b>						



# Recovery With Undo Logging

Examine each log entry  $\langle T, X, v \rangle$

- a) If T complete, do nothing.
- b) If T is incomplete, restore the old value of X

In what order?

**From most recent to earliest.**

# Nonquiescent Checkpoint (NQ CKPT)

- **Solution:** Let  $(T_1, \dots, T_k)$  be the active transactions
  1. Write  $\langle \text{START CKPT}(T_1, \dots, T_k) \rangle$   
and flush log to disk.
  2. Wait until all  $T_1 \dots T_k$  commit or abort,  
but don't prohibit new transactions.
  3. When all  $T_1 \dots T_k$  are “done”, write  $\langle \text{END CKPT} \rangle$  and flush  
log to disk.

# Recovery with NQ CKPT

## First case:

If the crash follows **<END CKPT>**

Then, undo

*any incomplete transaction that  
started after **<START CKPT>***

## Second case:

If the crash occurs between **<START CKPT>** and **<END CKPT>**

Then, undo

*any incomplete transaction that  
is on the **CKPT** list or  
started after **<START CKPT>***

# Undo Drawback

- Can't commit a transaction without first writing all its changed data to disk.
- Sometimes we can save disk I/O if we let changes to the DB reside only in main memory for a while;
- ...as long as we can fix things up in the event of a crash...

# Redo Logging

log all “important actions”

**<START T>** -- transaction **T** started.

**<T,X,New<sub>x</sub>>** -- database element **X** was modified;  
the new value is **New<sub>x</sub>**

**<COMMIT T>** -- transaction **T** has completed

**<ABORT T>** -- Transaction **T** couldn't complete successfully.

**One rule:**

R1. Before outputting **X** to disk, all log entries (including **<COMMIT T>**) must be written to log in disk.

# Example:

Action	t	Buff A	Buff B	A in HD	B in HD	Log
Read(A,t)	8	8		8	8	<Start T>
t:=t*2	16	8		8	8	
Write(A,t)	16	16		8	8	<T,A,16>
Read(B,t)	8	16	8	8	8	
t:=t*2	16	16	8	8	8	
Write(B,t)	16	16	16	8	8	<T,B,16> <Commit T>

## Flush Log

Output(A)	16	16	16	16	8
Output(B)	16	16	16	16	16

# Recovery With Redo Logging

Only committed transactions matter!

Examine each log entry  $\langle T, X, v \rangle$

a) If T incomplete, do nothing.

b) If T is complete, redo the operation:

For each  $\langle T, X, v \rangle$  in the log do:

`WRITE(X,v); OUTPUT(X);`

In what order?

**From the earliest to latest.**

# Checkpointing for Redo

1. Write a  $\langle \text{START CKPT}(T_1, \dots, T_k) \rangle$  record to the log, where  $T_i$ 's are all the active transactions.
2. Write to disk all the *dirty buffers* of transactions that had already committed when the **START CKPT** was written to log.
3. Write an  $\langle \text{END CKPT} \rangle$  record to log.



# Recovery with Ckpt. Redo

Two cases:

1. If the crash follows **<END CKPT>**,  
we can **restrict** ourselves to transactions that began after  
**<START CKPT>** and those in the **START** list.  
This is because we know that, in this case, every value written by  
committed transactions, before **<START CKPT(...)>**, is now in disk.
2. If the crash occurs between **<START CKPT>** and **<END CKPT>**,  
then go and find the previous **<END CKPT>** and do the same as in  
the first case.

# Functional Dependencies

# Superkeys and Keys

- $K$  is a *superkey* for relation  $R$  if  $K$  functionally determines all of  $R$ 's attributes.
- $K$  is a *key* for  $R$  if
  - $K$  is a superkey,
  - but no proper subset of  $K$  is a superkey.

# Boyce-Codd Normal Form

- **Boyce-Codd Normal Form (BCNF)**: simple condition under which the anomalies can be guaranteed not to exist.
- A relation R is in **BCNF** if:
  - Whenever there is a nontrivial dependency  
 $A_1 \dots A_n \rightarrow B_1 \dots B_m$   
for R, it must be the case that  
 $\{A_1, \dots, A_n\}$  is a **superkey** for R.

**Babies** isn't in **BCNF**.

- FD: **baby**  $\rightarrow$  **mother**
- Left side isn't a superkey.
  - We know: **baby** doesn't functionally determine **nurse**.

# Decomposition into BCNF

- One is **all the attributes** involved in the **violating dependency**
- The other is the **left side** and **all the other attributes** not involved in the dependency.
- By **repeatedly**, choosing suitable decompositions, we can break any relation schema into a collection of smaller schemas in BCNF.

# Babies Example

Births(baby, mother, nurse, doctor)

baby→mother is a violating FD, so we decompose.

Baby	Mother
Ben	Mary
Jason	Mary

Baby	Nurse	Doctor
Ben	Ann	Brown
Ben	Alice	Brown
Ben	Paula	Brown
Jason	Angela	Smith
Jason	Peggy	Smith
Jason	Rita	Smith

This relation needs to be further decomposed using the  
baby→doctor FD.

We, will see a formal algorithm for deducing this FD.

# Rules About Functional Dependencies

- Suppose **we are told** of a set of functional dependencies that a relation satisfies.

Based on them we can **deduce** other dependencies.

## Example.

$\text{baby} \rightarrow \text{mother}$  and

$\text{baby mother} \rightarrow \text{doctor}$

imply

$\text{baby} \rightarrow \text{doctor}$

But, what's  
the algorithm?

**Algorithm:** Starting with a set of attributes, repeatedly expand the set by adding the right sides of **FD**'s as soon as we have included their left sides.

If  $B \in \{A_1, A_2, \dots, A_n\}^+$  then FD:  $A_1 A_2 \dots A_n \rightarrow B$  **holds**.

If  $B \notin \{A_1, A_2, \dots, A_n\}^+$  then FD:  $A_1 A_2 \dots A_n \rightarrow B$  **doesn't hold**.

$\{A_1, A_2, \dots, A_n\}$  is a **superkey** iff  $\{A_1, A_2, \dots, A_n\}^+$  is the set of **all** attributes.

# Movie Example

Movies(title, year, studioName, president, presAddr)

and FDs:

title year  $\rightarrow$  studioName

studioName  $\rightarrow$  president

president  $\rightarrow$  presAddr

Last two violate **BCNF**. Why?

Compute  $\{\text{title, year}\}^+$ ,  $\{\text{studioName}\}^+$ ,  $\{\text{president}\}^+$  and see if you get all the attributes of the relation.

If not, you got a BCNF violation, and need to decompose.



# Example (Continued)

Let's decompose starting with:

studioName → president

Optional **rule of thumb**:

Add to the right-hand side any other attributes in the closure of studioName.

{studioName}<sup>+</sup> = {studioName, president, presAddr}

Thus, we get:

studioName → president presAddr

# Example (Continued)

Using:  $\text{studioName} \rightarrow \text{president}$   $\text{presAddr}$  we decompose into:

$\text{Movies1}(\text{studioName}, \text{president}, \text{presAddr})$

$\text{Movies2}(\text{title}, \text{year}, \text{studioName})$

$\text{Movie2}$  is in **BCNF**.

What about  $\text{Movie1}$ ?

FD  $\text{president} \rightarrow \text{presAddr}$  violates **BCNF**.

Why is it bad to leave  $\text{Movies1}$  as is?

*If many studios share the same president than we would have redundancy when repeating the **presAddr** for all those studios.*

# Example (Continued)

We decompose **Movies1**, using FD: **president**→**presAddr**

The resulting relation schemas, both in **BCNF**, are:

**Movies11**(**president**, **presAddr**)

**Movies12**(**studioName**, **president**)

So, finally we got **Movies11**, **Movies12**, and **Movies2**.

In general, we must keep applying the decomposition rule as many times as needed, until all our relations are in **BCNF**.

# Data Analysis with SQL Window Functions

# Table omc

<b>year</b>	<b>month</b>	<b>cat</b>	<b>countorders</b>	<b>revenue</b>
2009	10	ARTWORK	1782	45416
2009	10	BOOK	731	15299
2009	10	OCCASION	169	3476
2009	11	ARTWORK	2138	79390
2009	11	BOOK	2353	45808
2009	11	OCCASION	485	10041
2009	12	APPAREL	17	719
...	...	...	...	...

For each category: **which months were revenues below the average of the current year?**

```
SELECT year, month, cat, revenue,  
       AVG(revenue) OVER (PARTITION BY year, cat)  
AS avgrev  
FROM omc  
ORDER BY cat, year, month;
```

**Detail**

**Window**

In this example: all the tuples with the same **year** and **cat** as in the detail part.

## Second: Extract what you want with enclosing query

SELECT \*

FROM

(SELECT year, month, cat, revenue,

**AVG(revenue) OVER (PARTITION BY year,cat) AS avgrev**

FROM omc)

WHERE revenue < avgrev

ORDER BY cat, year, month;

Which months did the revenues from a product category drop below those of the same month of the previous year?

```
SELECT *
```

```
FROM (
```

```
    SELECT year, month, cat, revenue,
```

```
        LAG(revenue,12) OVER (PARTITION BY cat
```

```
                                ORDER BY year, month)
```

```
                                AS prev_year_rev
```

```
    FROM omc )
```

```
WHERE revenue < prev_year_rev
```

```
ORDER BY 1,2;
```



# Which are the top 10 months in terms of revenue for each category?

```
SELECT *  
FROM (  
  SELECT cat, year, month, round(revenue,-3),  
         RANK() OVER (PARTITION BY cat  
                     ORDER BY round(revenue,-3) DESC) AS rank  
  FROM omc  
)  
WHERE rank<=10  
ORDER BY cat, rank;
```

Similar to ROW\_NUMBER, but ties are not broken.  
See next slides for results.

If there were no ties,  
ROW\_NUMBER, RANK and  
DENSE\_RANK would be the same.